

### Question 1)

(a) You discover that your colleague wanted to target the general population of Taiwanese users of the product. However, he only collected data from a pool of young consumers, and missed many older customers who you suspect might use the product much less every day.

- Would this scenario create systematic or random error (or both or neither)?

→ systematic error

- Which part of the t-statistic or significance (diff, sd, n, alpha) would be affected?

- diff (yes, ↑)
- sd (no)
- n (no)
- alpha (no)

- Will it increase or decrease our power to reject the null hypothesis?

decrease power

- Which kind of error (Type I or Type II) becomes more likely because of this scenario?

Type II error

(b) You find that 20 of the respondents are reporting data from the wrong wearable device, so they should be removed from the data. These 20 people are just like the others in every other respect.

- Would this scenario create systematic or random error (or both or neither)?

→ random error

- Which part of the t-statistic or significance (diff, sd, n, alpha) would be affected?

- diff (no)
- sd (no)
- n (yes, ↓)
- alpha (no)

- Will it increase or decrease our power to reject the null hypothesis?

decrease power

- Which kind of error (Type I or Type II) becomes more likely because of this scenario?

Type II error

**(c) A very annoying professor visiting your company has criticized your colleague's "95% confidence" criteria, and has suggested relaxing it to just 90%.**

- Would this scenario create systematic or random error (or both or neither)?

→ systematic error

- Which part of the t-statistic or significance (diff, sd, n, alpha) would be affected?

- diff (no)
- sd (no)
- n (no)
- alpha (yes, ↑)

- Will it increase or decrease our power to reject the null hypothesis?

increase power

- Which kind of error (Type I or Type II) becomes more likely because of this scenario?

Type I error

**(d) Your colleague has measured usage times on five weekdays and taken a daily average. But you feel this will underreport usage for younger people who are very active on weekends, whereas it over-reports usage of older users.**

- Would this scenario create systematic or random error (or both or neither)?

→ systematic error

- Which part of the t-statistic or significance (diff, sd, n, alpha) would be affected?

- diff (no)
- sd (yes, ↑)
- n (no)
- alpha (no)

- Will it increase or decrease our power to reject the null hypothesis?

increase power

- Which kind of error (Type I or Type II) becomes more likely because of this scenario?

Type II error

## Question 2)

Let's return to the strictly fictional scenario (but with real data) from last week's Verizon dataset. Imagine this time that Verizon claims that they take no more than 7.6 minutes on average (single-tail test) to repair phone services for its customers. The file `verizon.csv` has a recent sample of repair times collected by the New York Public Utilities Commission, who seeks to verify this claim at 99% confidence.

(a) Recreate the traditional hypothesis test of last week using high-level built-in functions of R: (you may have to see the R help documentation, google how to use them, or ask for help on Teams)

```
repairtimes <- read.csv("verizon.csv", header=TRUE)$Time
```

(i) Use the `t.test()` function to conduct a one-sample, one-tailed t-test: report 99% confidence interval of the mean, t-value, and p-value

```
t.test(repairtimes, mu=7.6, conf.level = 0.99, alternative = "less")
```

```
##
## One Sample t-test
##
## data: repairtimes
## t = 2.5608, df = 1686, p-value = 0.9947
## alternative hypothesis: true mean is less than 7.6
## 99 percent confidence interval:
##      -Inf 9.360414
## sample estimates:
## mean of x
## 8.522009
```

(ii) Use the `power.t.test()` function to tell us the power of the test

```
hyp_mean <- 7.6
repairtimes_mean <- mean(repairtimes)
repairtimes_sd <- sd(repairtimes)
power.t.test(n = length(repairtimes), delta = repairtimes_mean-hyp_mean,
sd = repairtimes_sd, alternative = 'one.sided')
```

```
##
## Two-sample t test power calculation
##
##          n = 1687
##        delta = 0.9220095
##          sd = 14.78848
##    sig.level = 0.05
##        power = 0.5657309
## alternative = one.sided
##
## NOTE: n is number in *each* group
```

(b) Let's use bootstrapped hypothesis testing to re-examine this problem:

(i) Retrieve the original t-value from traditional methods (above)

```
repairtimes_se <- sd(repairtimes)/sqrt(length(repairtimes))
repairtimes_se
```

```
## [1] 0.3600527
```

```
t_value <- (repairtimes_mean-hyp_mean)/repairtimes_se
t_value
```


```
## [1] 2.560762
```

(ii) Bootstrap the null and alternative t-distributions

```
bootstrapped_null_alt <- function(sample0, hyp_mean) {
  resample <- sample(sample0, length(sample0), replace=TRUE)
  resample_se <- sd(sample0)/sqrt(length(resample))

  t_stat_alt <- (mean(resample)-hyp_mean)/resample_se
  t_stat_null <- (mean(resample)-mean(sample0))/resample_se

  c(t_stat_alt, t_stat_null)
}
boot_t_stats <- replicate(10000, bootstrapped_null_alt(repairtimes, hyp_mean))
t_alt <- boot_t_stats[1,]
t_null <- boot_t_stats[2,]
plot(density(t_alt), col="cornflowerblue", xlim=c(-4,6))
lines(density(t_null), lty="dashed")
abline(v=t_value, col="cornflowerblue")
```



(iii) Find the 99% cutoff value for critical null values of  $t$  (from the bootstrapped null); What should our test conclude when comparing the original  $t$ -value to the 99% cutoff value?

```
cutoff_99 <- quantile(t_null, probs=c(0.005, 0.995))
cutoff_99
```

```
##      0.5%      99.5%
## -2.449602  2.669968
```

(iv) Compute the p-value and power of our bootstrapped test

```
null_probs <- ecdf(t_null)
one_tailed_pvalue <- 1 - null_probs(t_value)
one_tailed_pvalue
```

```
## [1] 0.0067
```

```
alt_probs <- ecdf(t_alt)
alt_probs(cutoff_99[1]) + (1-alt_probs(cutoff_99[2]))
```

```
## [1] 0.4497
```